[22] Ludwig Schmidt, Shibani Santurkar, Dimitris Tsipras, Kunal Talwar, and Aleksander Madry. Adversarially robust generalization requires more data. In *Advances in Neural Information Processing Systems*, 2017.

[23] Ali Shafahi, W Ronny Huang, Christoph Studer, Soheil Feizi, and Tom Goldstein. Are adversarial examples inevitable? In *International Conference on Learning Representations*, 2019.

[24] Mahmood Sharif, Lujo Bauer, and Michael K Reiter. *n*-ML: Mitigating adversarial examples via ensembles of topologically manipulated classifiers. *arXiv preprint arXiv:1912.09059*, 2019.

[25] Christian Szegedy, Wojciech Zaremba, Ilya Sutskever, Joan Bruna, Dumitru Erhan, Ian Goodfellow, and Rob Fergus. Intriguing properties of neural networks. In *International Conference on Learning Representations*, 2014.

[26] Vincent Tjeng, Kai Xiao, and Russ Tedrake. Evaluating robustness of neural networks with Mixed Integer Programming. In *International Conference on Learning Representations*, 2019.

[27] Florian Tramer, Nicholas Carlini, Wieland Brendel, and Aleksander Madry. On adaptive attacks to adversarial example defenses. *arXiv:2002.08347*, 2020.

[28] Florian Tramèr, Alexey Kurakin, Nicolas Papernot, Ian Goodfellow, Dan Boneh, and Patrick McDaniel. Ensemble adversarial training: Attacks and defenses. In *International Conference on Learning Representations*, 2018.

[29] Eric Wong and Zico Kolter. Provable defenses against adversarial examples via the convex outer adversarial polytope. In *International Conference on Machine Learning*, 2018.

[30] Cihang Xie, Jianyu Wang, Zhishuai Zhang, Zhou Ren, and Alan Yuille. Mitigating adversarial effects through randomization. In *International Confernce on Learning Representations*, 2018.

[31] Weilin Xu, David Evans, and Yanjun Qi. Feature Squeezing: Detecting adversarial examples in deep neural networks. In *Network and Distributed Systems Security Symposium*, 2018.

[32] Xiao Zhang and David Evans. Cost-sensitive robustness against adversarial examples. In *International Conference on Learning Representations*, 2019.

A Additional Experimental Results

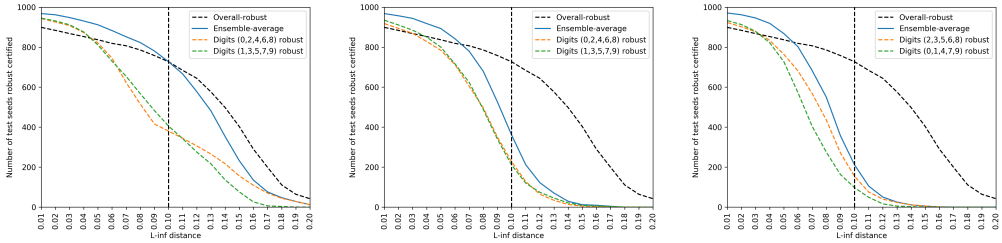

(a) Even/odd seeds robust (b) Even/odd targets robust (c) Clustered Targets

Figure 6: Number of test seeds certified to be jointly robust using the individual models and different two-model ensembles average framework for different ϵ values.

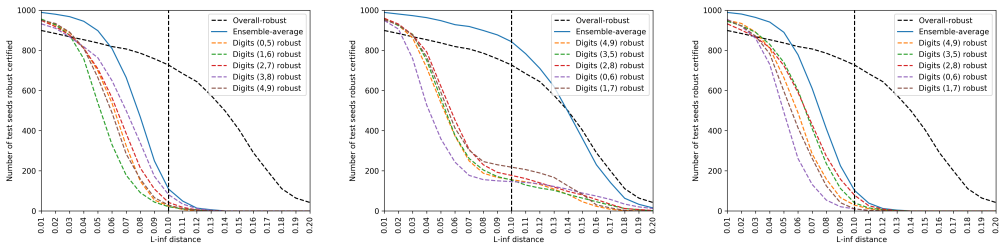

(a) Modulo-5 Targets Robust (b) Modulo-5 Seeds Robust (c) Clustered Targets

Figure 7: Number of test examples certified to be jointly robust using the individual models and the 5-model ensembles average framework with different ϵ values.

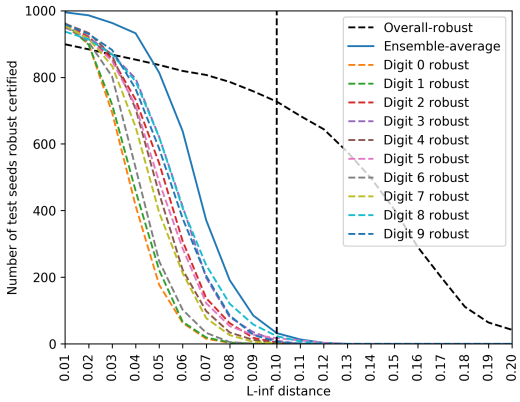

Figure 8: Number of test examples certified to be jointly robust using the individual models and the 10-model ensemble average framework with targets robust for different ϵ values.